VIZIT

RESEARCH REPORT
2026

THE STATE OF AGENTIC SHOPPING

Agentic shopping: how AI agents *discover* and *recommend* products

Why your product images are the most important signal you are not optimizing.

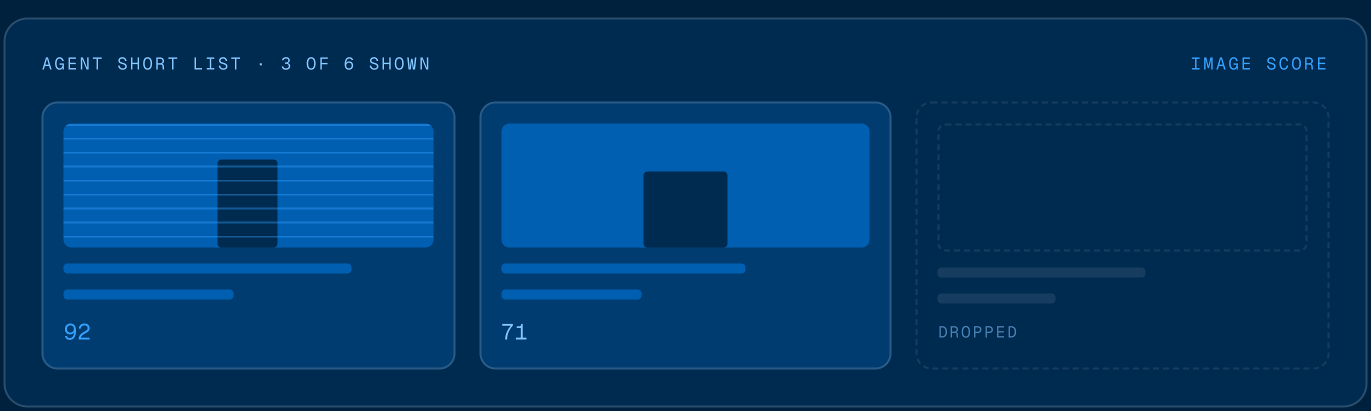

CONTENTS

What is *inside*

The shift is *here*

A year ago, AI-mediated product discovery was a novelty. Today it is a primary shopping channel.

The numbers tell a clear story. AI-driven traffic to retail sites grew **393% year over year**, and during the 2025 holiday season that figure hit 693%.[1] Shopify merchants alone saw AI-referred traffic climb roughly 8x, with AI-driven orders surging 15x.[2] Nearly **half of all shoppers (45%) now use AI somewhere in their buying journey**, according to IBM's 2026 Commerce Index.[3] And 58% say they have replaced traditional search with generative AI for product recommendations.[3]

This is not experimental behavior from early adopters. It is mass-market adoption happening faster than mobile commerce did a decade ago.

What makes this shift different from previous channel disruptions is the mechanics. AI agents do not scroll through pages of results the way a human does. They retrieve, rank, and present a short list of three to six products. If your product is not on that list, there is no “page two” to fall back on. You are invisible.

The agents doing this work operate across three distinct surfaces: inside retailer ecosystems, inside general-purpose AI tools, and inside search engines. Each surface works differently. Each presents different risks and opportunities. And all three are growing simultaneously.

Here is the part that should get your attention: **AI-referred shoppers convert 31% higher than traffic from other sources**.[4] They are 39% more likely to try new brands and spend 37% more per order.[5] These are not window shoppers. They are high-intent buyers whose path to purchase is now mediated by an algorithm you probably are not optimizing for.

> The question is not whether AI agents will influence your sales. It is whether your product content is *ready* for how they work.

1, 4, 5 — Salesforce Commerce Cloud data, 2025-2026 reporting period. 2 — Shopify merchant analytics, 2025-2026 YoY. 3 — IBM 2026 Commerce Index.

Agentic shopping is not *arriving*. It has arrived.

393%

Year-over-year growth in AI-driven traffic to retail sites

693% OVER THE 2025 HOLIDAY SEASON[1]

45%

of shoppers now use AI somewhere in their buying journey

58% HAVE REPLACED SEARCH WITH GEN AI[3]

31%

higher conversion from AI-referred shoppers than other traffic

+37% ORDER VALUE · +39% BRAND TRIAL[4]

8x

growth in AI-referred traffic across Shopify merchants

15X GROWTH IN AI-DRIVEN ORDERS[2]

AI-REFERRED RETAIL TRAFFIC, INDEXED

INDEX 100 = 2024

1, 4 — Salesforce Commerce Cloud, 2025-2026. 2 — Shopify merchant analytics. 3 — IBM 2026 Commerce Index.

Where agents *operate*: three surfaces

Agentic shopping is not a single channel. It is three overlapping surfaces, each with different rules, different data access, and different implications for brands.

SURFACE ONE

Retailer-native agents

A shopper asks a retailer's assistant for a high-fiber cereal under $6. Four product cards come back with images, prices, and a line on why each fits. She taps one and buys, in under 90 seconds.

HOW IT WORKS These agents work inside a closed ecosystem, searching only that retailer's catalog with the retailer's own product data, imagery, and reviews. Amazon retired Rufus in early 2026 and launched Alexa for Shopping with a "Buy for Me" feature that completes purchases autonomously.[6]

FOR BRANDS Your listing is your storefront. Walmart reports Sparky users spend **35% more per order**, and roughly half of app users have tried it.[7]

SURFACE TWO

General-purpose agents

A shopper asks a chat assistant for the best lunchbox snack under $5 a pack. It returns five options with photos, prices, and a short analysis of each, pulled from multiple retailers and review sites.

HOW IT WORKS They search the open web, aggregating retailer listings, review sites, brand sites, and editorial content. Perplexity launched Buy with Pro and Snap to Shop for visual search.[8] ChatGPT paused Instant Checkout in March 2026 to focus on a discovery-first model.[9]

FOR BRANDS You can be compared and recommended (or excluded) before the shopper visits any retailer. A new consideration layer above retail search.

SURFACE THREE

Search-integrated agents

A shopper photographs a cereal box on a friend's counter and instantly gets visually similar products with prices and purchase links, before typing a single brand name.

HOW IT WORKS Google AI Mode now reaches over **1 billion monthly users**, and Google Lens processes roughly **20 billion searches per month**, growing about 140% year over year.[10] Matching runs on visual embeddings, so image clarity decides whether you appear.

FOR BRANDS This surface is the top of the funnel. Image quality, distinctiveness, and machine-readability are direct ranking factors.

6 — Amazon Alexa for Shopping launch, May 2026. 7 — Walmart Q1 2026 earnings call and Sparky usage disclosures. 8 — Perplexity product announcements, 2025-2026. 9 — OpenAI blog, March 2026. 10 — Google I/O 2026 and Google Search statistics.

All three surfaces run at once, and they feed each other.

RETAILER-NATIVE

amazon

Walmart

Images display in the result card and feed attribute extraction from your listing.

GENERAL-PURPOSE

ChatGPT

perplexity

Gemini

Images display in comparisons and are parsed alongside web-wide structured data.

SEARCH-INTEGRATED

Google AI

Google Lens

The image is the primary matching signal. Visual embeddings decide who appears.

ONE SHOPPER, THREE SURFACES

Discovers via visual search → Researches in a chat agent → Buys in a retailer app

	RETAILER-NATIVE	GENERAL-PURPOSE	SEARCH-INTEGRATED
CATALOG SCOPE	Single retailer	Cross-retailer, open web	Cross-retailer, visual index
PRIMARY INPUT	Retailer product data	Web-wide data + reviews	Visual embeddings + structured data
IMAGE ROLE	Display + attribute extraction	Display + structured parsing	Primary matching signal
SHOPPER INTENT	High — already on retailer	Research + comparison	Discovery — pre-retailer
BRAND ACTION	Optimize listing content	Cross-platform data quality	Optimize imagery for machine vision

MECHANICS

How agents *choose* products

Every AI shopping agent follows the same basic process, regardless of which surface it operates on. That compression, from thousands of candidates to a handful of recommendations, is the single most important dynamic in agentic shopping.

FIGURE 03 — THE AGENTIC FUNNEL

RETRIEVE → RANK → PRESENT

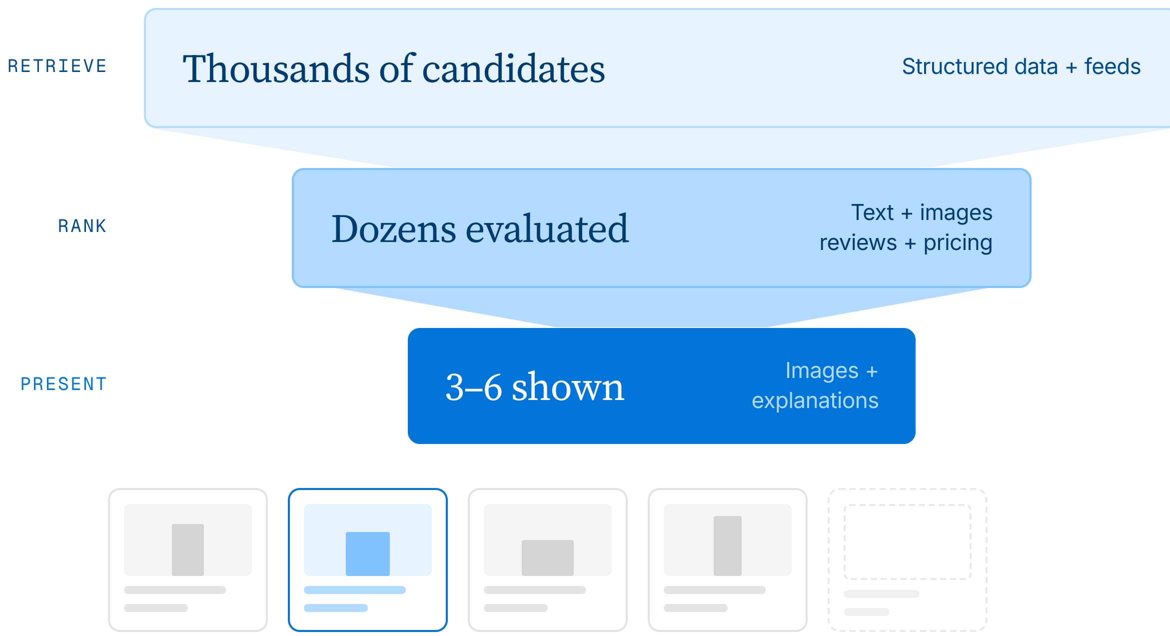

THE SHOPPER SEES ONLY THIS ROW. THERE IS NO PAGE TWO.

> You are either on the short list, or you *do not exist*.

SIGNALS

What agents actually evaluate

TABLE STAKES

Text and structured data

Titles, descriptions, attributes, specifications. Without clean structured data you never pass retrieval.

CONFIDENCE MULTIPLIER

Reviews and trust signals

Volume, recency, and rating together. 4.2 stars with 2,000 reviews outranks 4.5 stars with 12.

HARD GATE

Pricing and availability

Agents penalize products where price cannot be confidently extracted.

MOST UNDERINVESTED

Images

Not a display asset at the end of the funnel. An active input at every stage of it.

FIGURE 04 — THE PROTOCOL LAYER

Two open protocols are standardizing what agents can read

Both require structured product data with image fields. They are a forcing function: image quality becomes a measurable, indexable signal rather than a subjective design choice.

ACP

OPENAI + STRIPE

Agent Commerce Protocol. Powers ChatGPT Shopping, available as an open-source standard. Defines how agents access catalogs, including image fields, pricing, and availability.

+

UCP

GOOGLE + SHOPIFY

Unified Commerce Protocol. Announced at NRF in January 2026, expanded in May. Standardizes how product data, including visual assets, flows to agents.

≈40% more agentic traffic for merchants supporting **both** protocols, versus those supporting only one or neither.[11]

11 — Agentic commerce protocol adoption data, industry analysis 2026.

Five jobs your product images are *now* doing

The conversation about agentic readiness has focused almost entirely on text. And text matters. Clean structured data is how products get into the retrieval set. But once products are retrieved, images do the load-bearing work at multiple points in the funnel. Most brands are only thinking about the first of these five jobs.

01	**Display asset for a shorter, higher-stakes list** Every image is now a hero image. It must win the click from a list of five, with no scroll-back.	PRESENT STAGE
02	**Machine-readable data extraction** Agents read packaging with OCR and vision. Your label is now a data source, not just branding.	RETRIEVE + RANK
03	**Causally changing model conclusions** The image is evidence, not decoration. Mismatched imagery measurably lowers attribute accuracy.	RANK STAGE
04	**Confidence gating** Below a confidence threshold the agent silently drops you. There is no impressions metric for that.	RANK → PRESENT
05	**The image as query input** In visual search your image is the query. Distinctiveness and fidelity are ranking factors.	PRE-RETRIEVE

01 The display asset for a shorter, higher-stakes list

When an agent presents three to six products, the human's final selection is overwhelmingly visual: an image, a title, a price, maybe one sentence of reasoning. The image is doing more persuasion work per impression than it ever did on a crowded results page.

> BUSINESS IMPLICATION
> Every product image is now a hero image. If it does not win the click from a list of five, you have lost the sale.

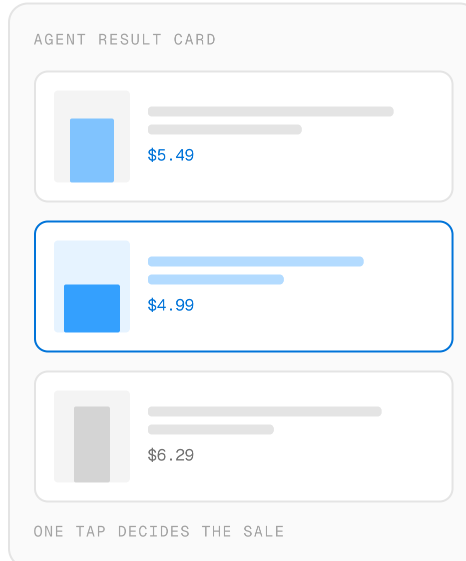

02 Machine-readable data extraction

Agents do not just display your images. They read them. Through OCR and visual attribute extraction they pull structured data straight off your packaging: ingredient lists, nutrition panels, size specifications, color variants. If it is visible in the image, an agent can extract it and use it for matching and ranking.

> BUSINESS IMPLICATION
> Your packaging is a data source, not just a branding vehicle. A cluttered, text-heavy hero image yields lower-confidence reads than a clean, well-lit product shot, and that directly affects match accuracy.

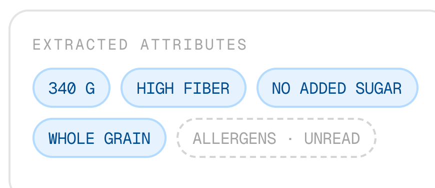

03 Causally changing model conclusions

This is the job most brands do not know about. Controlled studies show that **product images causally change the conclusions AI models reach about a product's attributes.**[12] Shown listings with correct text but low-quality images, models' attribute-prediction accuracy dropped by a statistically significant margin. The image is not decoration. It is evidence, and when image and text conflict, confidence falls.

> BUSINESS IMPLICATION
> An ambiguous image can make an agent draw the wrong conclusion about what your product is, what it does, and who it is for. That is a measurable effect on whether you get recommended.

12 — Controlled multimodal attribute prediction studies, 2025-2026.

04 Confidence gating

Agents have confidence thresholds. When an agent cannot confidently assess a product, it often chooses not to surface it at all rather than risk a bad recommendation. Low-resolution images, heavy text overlays, cluttered lifestyle shots, and composites with multiple products all reduce confidence in visual analysis.

BUSINESS IMPLICATION

A product with poor imagery may not just rank lower. **It may never appear at all.** It is dropped silently, with no impressions metric to tell you it happened.

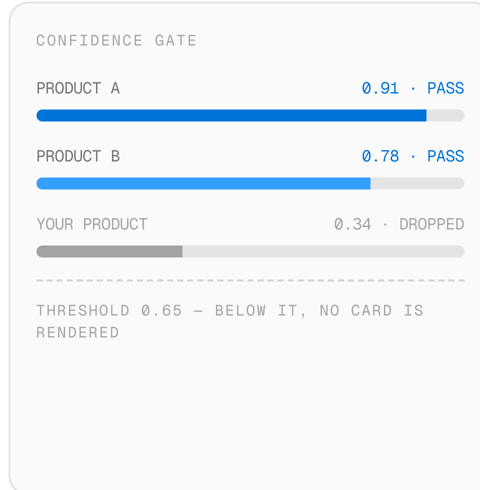

05 The image as query input

With Google Lens processing 20 billion visual searches a month, a growing share of discovery starts with an image, not a text query. A shopper photographs a product in the wild and the agent matches it against a visual index. Here your image is not describing the product. It *is* the thing being matched.

BUSINESS IMPLICATION

If your hero looks like every other pack in the category, visual search cannot distinguish you. Distinctiveness is a ranking factor.

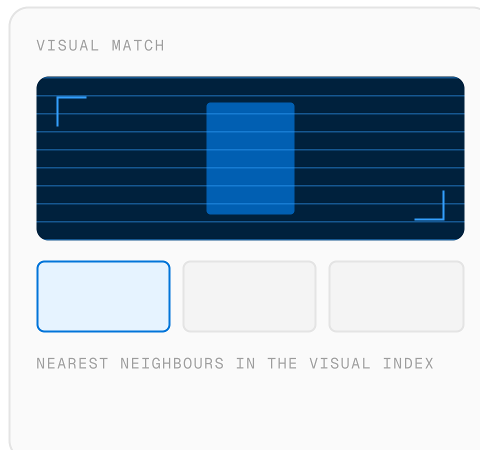

THE MIRAGE CAVEAT

Frontier models will confidently describe images they have never actually seen more than *60% of the time.*[13]

If you do not supply high-quality, accurate product imagery, agents will not simply skip your product. They may invent details about it. If you control the image input, you control the output. If you leave gaps, agents fill them with guesses.

13 — Frontier model hallucination benchmarks, AI research 2025-2026.

FIGURE 05 — HOW AN AGENT READS TWO HERO IMAGES

SAME PRODUCT

The same product, read two ways

CLUTTERED HERO

AS UPLOADED

AGENT CONFIDENCE 0.34

- Two products in frame: which one is listed?
- Overlays read as product text, polluting attributes
- Promotional badges misread as pack claims
- Busy background lowers segmentation quality

RESULT — DROPPED FROM THE SHORT LIST

CLEAN HERO

MACHINE-READY

SINGLE SUBJECT · CLEAR GROUND

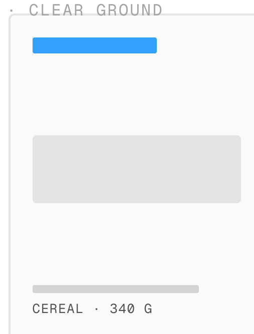

AGENT CONFIDENCE 0.91

- One subject, unambiguous product identity
- Pack copy extracts cleanly into attributes
- No promotional text to misinterpret
- Distinct silhouette for visual-search matching

RESULT — RECOMMENDED BY THE AGENT

> Text gets products into the consideration set. Images do the *load-bearing work* at every stage after that.

READINESS

Most brands are *not* ready

If images are this important, you would expect heavy investment in visual content optimization for AI agents. Readiness audits across CPG and retail brands say otherwise: **most score below 25 out of 60** on a composite agentic readiness assessment.[14] The weakest dimension, consistently, is image optimization for machine readability.

FIGURE 06 — AGENTIC READINESS, SIX DIMENSIONS

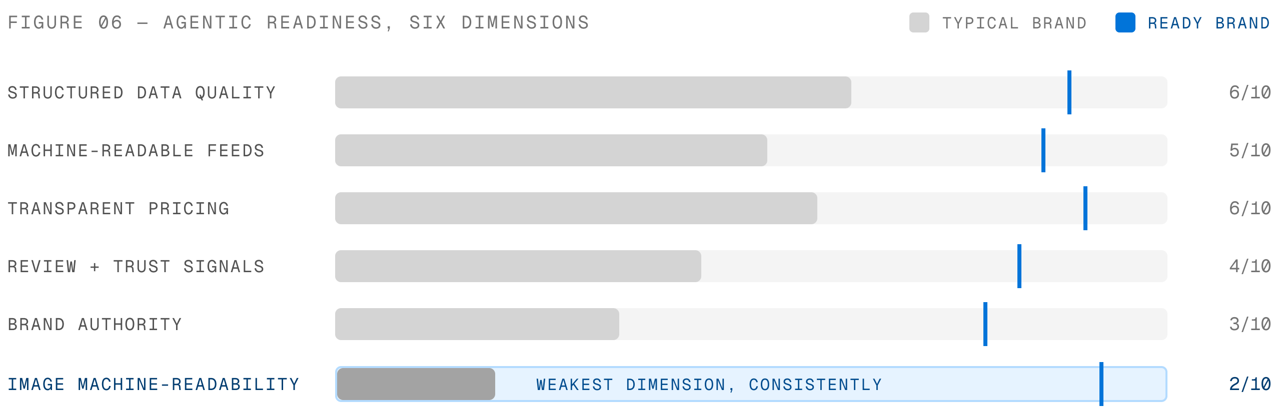

26/60

Composite score for a representative brand. Text-side dimensions carry the total; visual readiness drags it down, and visual readiness is what the ranking, gating, and final-click stages depend on.

THREE COMMON MISTAKES

Treating images as done at launch. Photography shot once for human browsers, never re-evaluated for how an agent parses it.

Heroes designed to impress, not inform. The promotional badge a designer added becomes noise that lowers the agent's confidence.

Optimizing text while ignoring visuals. Feed and keyword work address retrieval only. Images decide ranking, gating, and the final click.

> When an agent surfaces three to six products and yours is not among them, you get zero visibility for that query. Not reduced visibility. *Zero*.

14 — Agentic readiness audits across CPG and retail brands, 2026.

How to *win*

Agentic readiness is not a vague aspiration. It is a set of specific optimizations you can start this quarter.

Audit your product imagery for machine readability
01

Review every hero through the lens of a vision model, not a human shopper. Can it identify the product? Extract size, color, and key features? Are overlays confusing the read?

Make the hero image parseable
02

It does triple duty: display asset, data source, and visual signature. Minimize overlays, use a clean ground, make the product the focal point. Lifestyle and badges belong in secondary slots.

Build structured data completeness across both protocols
03

Support ACP and UCP if your stack allows. Dual-protocol merchants see roughly 40% more agentic traffic. Ship complete image fields, accurate attributes, transparent pricing.

Maintain review freshness
04

Agents treat volume, recency, and sentiment as confidence multipliers. Strong reviews from the past 90 days beat strong reviews from two years ago. Build generation into ongoing operations.

Measure how your products perform with AI agents
05

This is the gap most brands have not closed. You measure click-through, conversion, and ranking. Do you measure how agents perceive your imagery, or whether you appear in agentic short lists for your target queries?

> The defensible advantage in agentic shopping is not a secret. It is *discipline*.

HOW VIZIT HELPS

Built for *this* moment

Vizit has spent years helping global, household brands understand how their visual content performs. For most of that time the question was straightforward: which product images resonate best with human shoppers?

That question still matters. But a second audience now evaluates the same images: agents that retrieve, rank, and recommend across every major shopping surface. The visual signals that influence a shopper's click also decide whether an agent surfaces your product at all. What Vizit has always done now matters more than it ever has.

THE JOB YOUR IMAGE IS DOING	HOW VIZIT ADDRESSES IT
Winning the human's selection Three to six products, one visual choice.	**Image Score** Quantifies how an image performs with specific category shoppers, so you know which shot wins the click before you publish it.
Machine-readable extraction Agents read attributes off your pack.	**Visual clarity and composition analysis** The qualities that make an image easy for a human to parse at a glance make it easy for a machine to parse programmatically.
Shaping model conclusions Color, texture, context, composition.	**Category Context Models** Category-specific models let you test how different images shape conversion, connecting visual drivers and sales.
Staying above the threshold Silent drops you never see.	**Underperformance diagnostics** Identifies which images in your catalog are underperforming and why. It is the visibility most brands have never had.
Standing out in visual search Distinctiveness is the ranking signal.	**Competitive distinctiveness analysis** Shows how your imagery compares to your category, and whether you can be matched or blend into a sea of similar packs.
THE FULL PICTURE	**Product Page Score** Measures how well a product listing converts shoppers relative to category competitors.

HOW VIZIT HELPS

Two ways to measure what agents and shoppers see

PRODUCT

Image Score

92 / 100

Quantifies how well an image performs with specific target audiences and AI agents. No more guessing which image looks better. Vizit tells you which one performs better, and why.

PRODUCT

Product Page Score

- IMAGERY
- COPY
- DATA

Evaluates how an entire listing performs, combining visual content analysis with the structured data signals that determine whether agents surface your product.

What has changed is the *stakes*, not the work.

The disciplines Vizit has always championed (testing visual content, measuring performance by audience, optimizing on data instead of opinion) are the disciplines the agentic era demands. The difference is that the cost of getting it wrong has gone from a lower conversion rate to invisibility across an entire shopping channel.

A FREE PRODUCT PAGE SCORECARD

Get a free *audit* of your product listing.

Request a free audit

SOURCES

01 Salesforce Commerce Cloud data, 2025–2026 reporting period

02 Shopify merchant analytics, 2025–2026 YoY comparison

03 IBM 2026 Commerce Index

04 Salesforce Commerce Cloud conversion analysis, 2025–2026

05 Salesforce Commerce Cloud shopper behavior analysis, 2025–2026

06 Amazon Alexa for Shopping launch announcement, May 2026

07 Walmart Q1 2026 earnings call and Sparky usage disclosures

08 Perplexity product announcements, 2025–2026

09 ChatGPT commerce feature timeline, OpenAI blog, March 2026

10 Google I/O 2026 and Google Search statistics

11 Agentic commerce protocol adoption data, industry analysis 2026

12 Controlled multimodal attribute prediction studies, 2025–2026

13 Frontier model hallucination benchmarks, AI research 2025–2026

14 Agentic readiness audits across CPG and retail brands, 2026

VIZIT

The AI Product Conversion Platform

VIZIT.COM